

# CSC 470 – Section 3

## Topics in Computer Science: Advanced Browser Technologies

Mark F. Russo, Ph.D.

Spring 2016

Lecture 12: Image Processing with Web Workers


# Color Models

Pixel colors are represented using different models

- Popular color models decompose a single color into three orthogonal components
- Dimensionality of the "color space" varies depending upon the model used

Popular Color Models Include:

- RGB – Red, Green, Blue
- HSV – Hue, Saturation, Value
- HSL – Hue, Saturation, Lightness



# Color Models - RGB

Most popular is Red-Green-Blue (RGB)

Cartesian coordinate system

Each of three dimensions represents a pure color

- Red, Green, and Blue

The value of each dimension is stored in one byte

- How large of a value can each dimension take on?

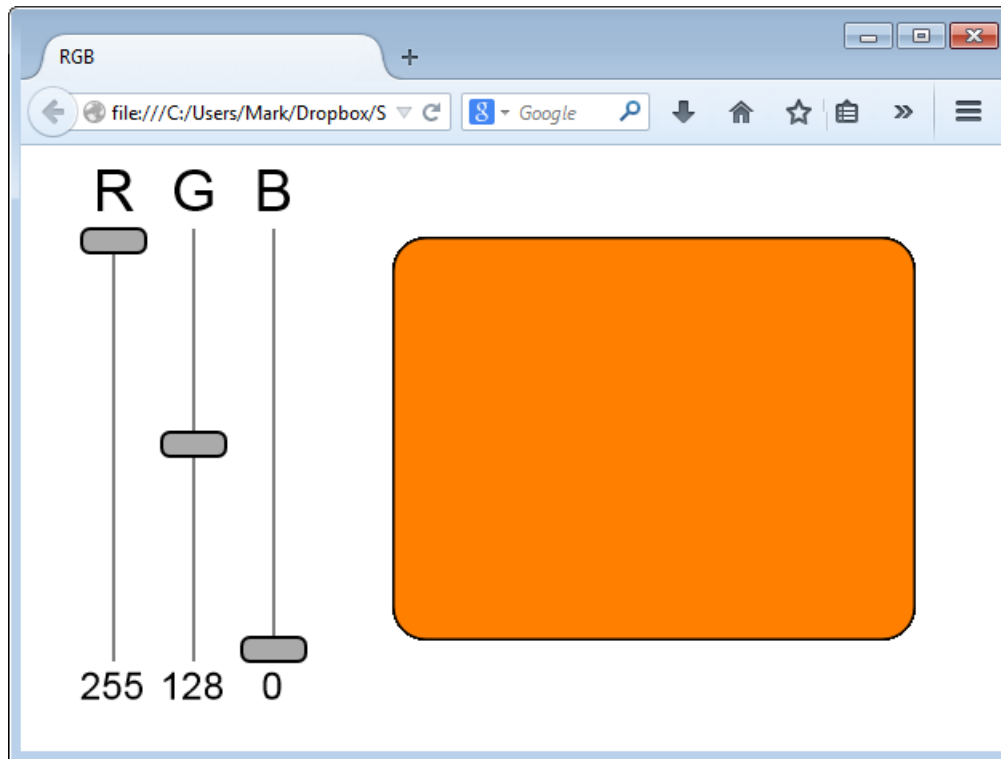
Origin (0, 0, 0) is black

On-axis points are shades of pure colors

Off-axis points represent other colors

(255, 255, 255) is white

# Color Models - RGB



Find

- pure red
- pure blue
- pure green
- black
- white
- gray
- yellow
- magenta
- cyan

# Color Models - HSV

## Hue-Saturation-Value (HSV)

- Cylindrical coordinate system
- **Hue** - Pure color value from rainbow (angle)
- **Saturation** - Grayscale to pure color (radius)
- **Value** - Black to color (height)


Useful when intensity or hue is required independent of the other

- Eg. Color contour plots

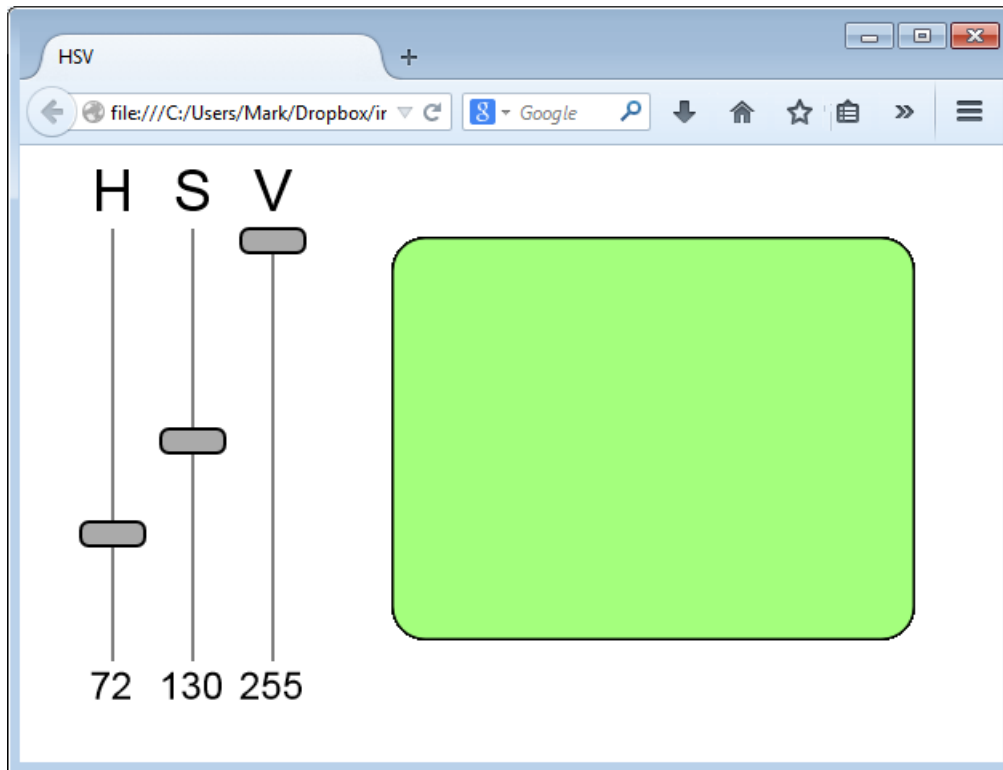
Black: (any, any, 0)

White: (any, 0, 255)

Grayscale: (any, 0, any)



# Color Models - HSV




Find

- pure red
- pure blue
- pure green
- black
- white
- gray
- yellow
- magenta
- cyan


# Transparency

- The degree to which a pixel allows the color of the pixel "below it" to bleed through



# 32-bit RGBA Color Model

- Adds a fourth value (called alpha) to a pixel color triple
  - alpha = 255                      fully opaque
  - alpha = 0                        fully transparent
- If three pixel color values and alpha are each stored in one byte, a pixel will be 32-bits (4 bytes) in total
- **RGBA** color byte example:





# Image File Formats

- An image file can have a variety of standard formats
- Image files include information such as:
  - **Pixel data** : pixel color and location
  - **Header information** : non-pixel data
    - date-time image captures
    - camera settings
    - GPS coordinates at which picture was taken
    - ...

# Image Files - Compression

- Algorithms used to encode information in fewer bytes than the original
- Used in image files to reduce file size/ required disk storage
- May or may not result in data loss

## **File size depends upon:**

- image dimensions
- color-depth
- compression algorithm


# "Deflate" Compression

- Used by Zip archives and PNG image files
- Repeated byte patterns are identified in data stream
- When found, a repeat is replaced by the location in the stream it was last seen along with its length
- Patterns are 4-258 bytes in length
- Patterns are further encoded as a (Huffman) tree

# "Deflate" Compression

*Mary had a little lamb, little lamb, little lamb.*

*Mary had a little lamb, its fleece was white as snow.*



The diagram illustrates the Deflate compression process. It shows three red arrows originating from the text above and pointing to specific locations in the compressed text below. The first arrow starts at the beginning of the first line and points to the start of the second line. The second arrow starts at the beginning of the second line and points to the start of the third line. The third arrow starts at the beginning of the third line and points to the start of the fourth line. This represents the sequence of jumps and replays in the compressed data.

*Mary had a little lamb, (•,13)(•,11). (•,24) its fleece  
was white as snow.*

*(•,13)  $\equiv$  Jump to identified location and replay 13 characters*

# Image Files – Lossy vs. Lossless

## **Lossless File Formats**

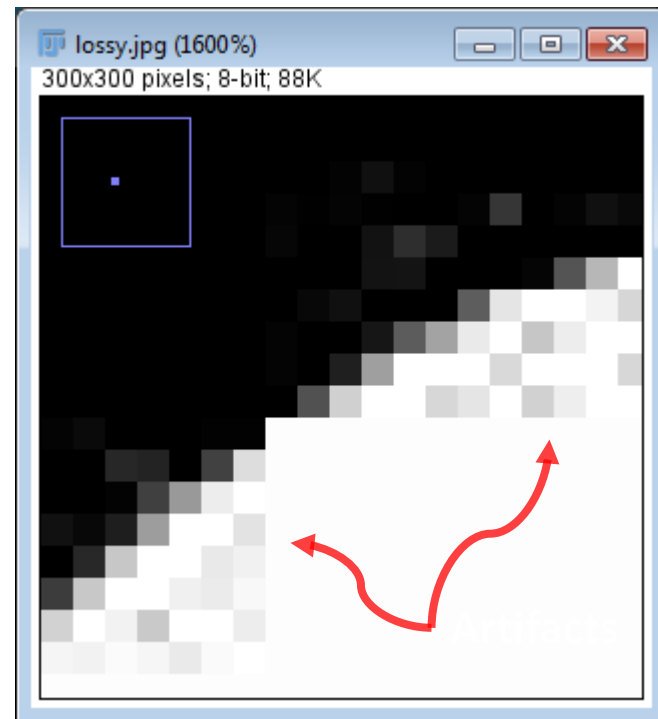

- No pixel information is lost, in spite of compression
- Image can be restored exactly as captured
- Ex: TIFF, PNG, GIF, BMP, ...

## **Lossy File Formats**

- Compression algorithm approximates pixel values
- Image cannot be restored to original
- Greater compression ratio = smaller files
- Ex: JPEG


# Lossy vs. Lossless

Artifacts in the JPEG file, which are introduced by the lossy compression algorithm



# Sampling Theorem

*If a function  $x(t)$  contains no frequencies higher than  $W$  cycles-per-second (hertz), it is completely determined by sampling at a series of points spaced  $1/2W$  seconds apart.*



*Aliasing occurs when different signals are indistinguishable due to inadequate sampling rate*

- For adequate spatial resolution of image features there must be at least 2 pixels per resolvable unit
- Credited to Harry Nyquist and Claude Shannon

# Moiré Pattern

- When the Sampling Theorem's pixel density requirements are not satisfied a *Moiré Pattern* can result - a kind of aliasing



Downsampling  
→





# Image Processing

Three broad algorithm groups:

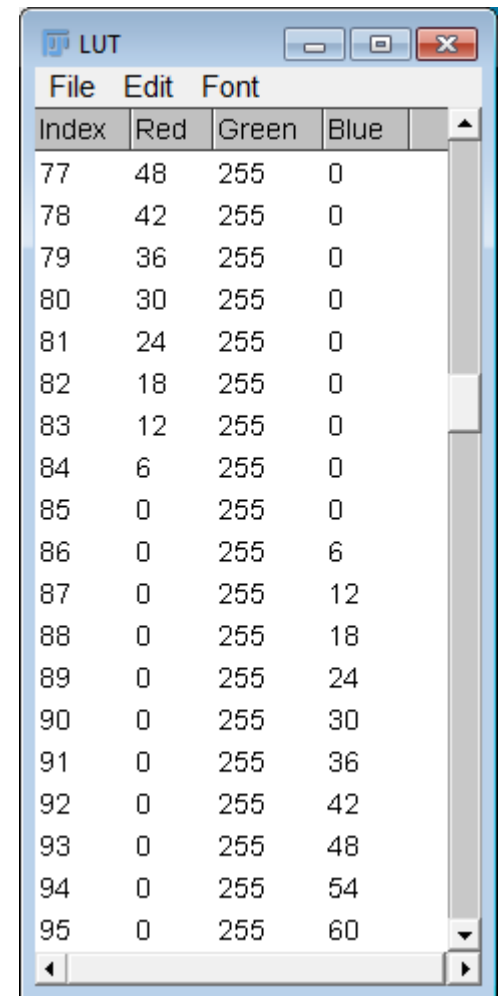
1. Intensity Transformation Filters – Point Operations
  - Output image pixel color depends on corresponding input pixel color only
2. Local Filtering/Convolution – Pixel Region Operations
  - Kernel application
  - Output image pixel color depend on pixels near corresponding input pixel
3. Whole Image Algorithms
  - Algorithms that make use of the entire image to perform processing

# Point Operations

- Lookup Tables
- Invert
- Grayscale
- Sepia
- Brighten (add an offset to all pixels)
- Saturate
- Posterize
- Threshold
- Image Math (add images, subtract images, ...)

# Lookup Tables (LUTs)


- Maps 8-bit grayscale pixel values to RGB colors
- Used to transform an image by "looking up" gray scale values and replacing with corresponding color
- A LUT is chosen strategically to highlight features of interest
  - E.g. Use hue to highlight subtle grayscale features




A screenshot of a software window titled "LUT" with a menu bar containing "File", "Edit", and "Font". The window displays a table with four columns: "Index", "Red", "Green", and "Blue". The table contains data for indices 77 through 95. The "Red" column is constant at 48, and the "Green" column is constant at 255. The "Blue" column increases linearly from 0 at index 77 to 60 at index 95. The window has standard OS controls (minimize, maximize, close) in the top right corner and a scrollbar on the right side.


Index	Red	Green	Blue
77	48	255	0
78	42	255	0
79	36	255	0
80	30	255	0
81	24	255	0
82	18	255	0
83	12	255	0
84	6	255	0
85	0	255	0
86	0	255	6
87	0	255	12
88	0	255	18
89	0	255	24
90	0	255	30
91	0	255	36
92	0	255	42
93	0	255	48
94	0	255	54
95	0	255	60

# Predefined (Named) LUTs







edges



ICA



# Point Operations

## **grayscale:**

$r = g = b = 0.299 * r + 0.587 * g + 0.114 * b;$

## **sepia:**

$r = (r * 0.393) + (g * 0.769) + (b * 0.189);$

$g = (r * 0.349) + (g * 0.686) + (b * 0.168);$

$b = (r * 0.272) + (g * 0.534) + (b * 0.131);$

## **invert:**

$r = 255 - r;$

$g = 255 - g;$

$b = 255 - b;$

## **brighten (p):**

$r = \text{clamp}(r + p);$

$g = \text{clamp}(g + p);$

$b = \text{clamp}(b + p);$

## **saturate: (p in [0.0, 2.0])**

$\text{var avg} = (r + g + b) / 3;$

$r = \text{avg} + p * (r - \text{avg});$

$g = \text{avg} + p * (g - \text{avg});$

$b = \text{avg} + p * (b - \text{avg});$

## **posterize: (p in [1, 255])**

$\text{var step} = \text{Math.floor}(255 / p);$



$r = \text{Math.floor}(r / \text{step}) * \text{step};$

$g = \text{Math.floor}(g / \text{step}) * \text{step};$

$b = \text{Math.floor}(b / \text{step}) * \text{step};$



# Image Segmentation

- The process of partitioning a digital image into multiple segments of interest
- Example:
  - Highlight the "interesting" part of an alloy image in preparation for analysis




# Thresholding

- Simplest method of image segmentation
- Convert an image into a binary format
  - Parts of interest converted to one color
  - Everything else converted to another color
- Further processing on pixels with color of interest






# Thresholding



Pixel value histogram




# Local Filtering

## Examples:

- Image Convolution
  - Identity
  - Smooth (Average)
  - Box Blur
  - Gaussian Blur
  - Sharpen
  - Edge Detection
- Median Filter
- Dilation
- Erosion
- Opening
- Closing

# Local Filtering – Convolution Kernel



$$E' = w_1A + w_2B + w_3C + w_4D + w_5E + w_6F + w_7G + w_8H + w_9I$$

# Convolution – Identity Kernel

- No change

0	0	0
0	1	0
0	0	0


# Convolution – Box Blur

- Set pixel to the average of all colors in the neighborhood
- Smooths out areas of sharp changes.


$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

$1/25$	$1/25$	$1/25$	$1/25$	$1/25$
$1/25$	$1/25$	$1/25$	$1/25$	$1/25$
$1/25$	$1/25$	$1/25$	$1/25$	$1/25$
$1/25$	$1/25$	$1/25$	$1/25$	$1/25$
$1/25$	$1/25$	$1/25$	$1/25$	$1/25$


# Box Blur Example



Original

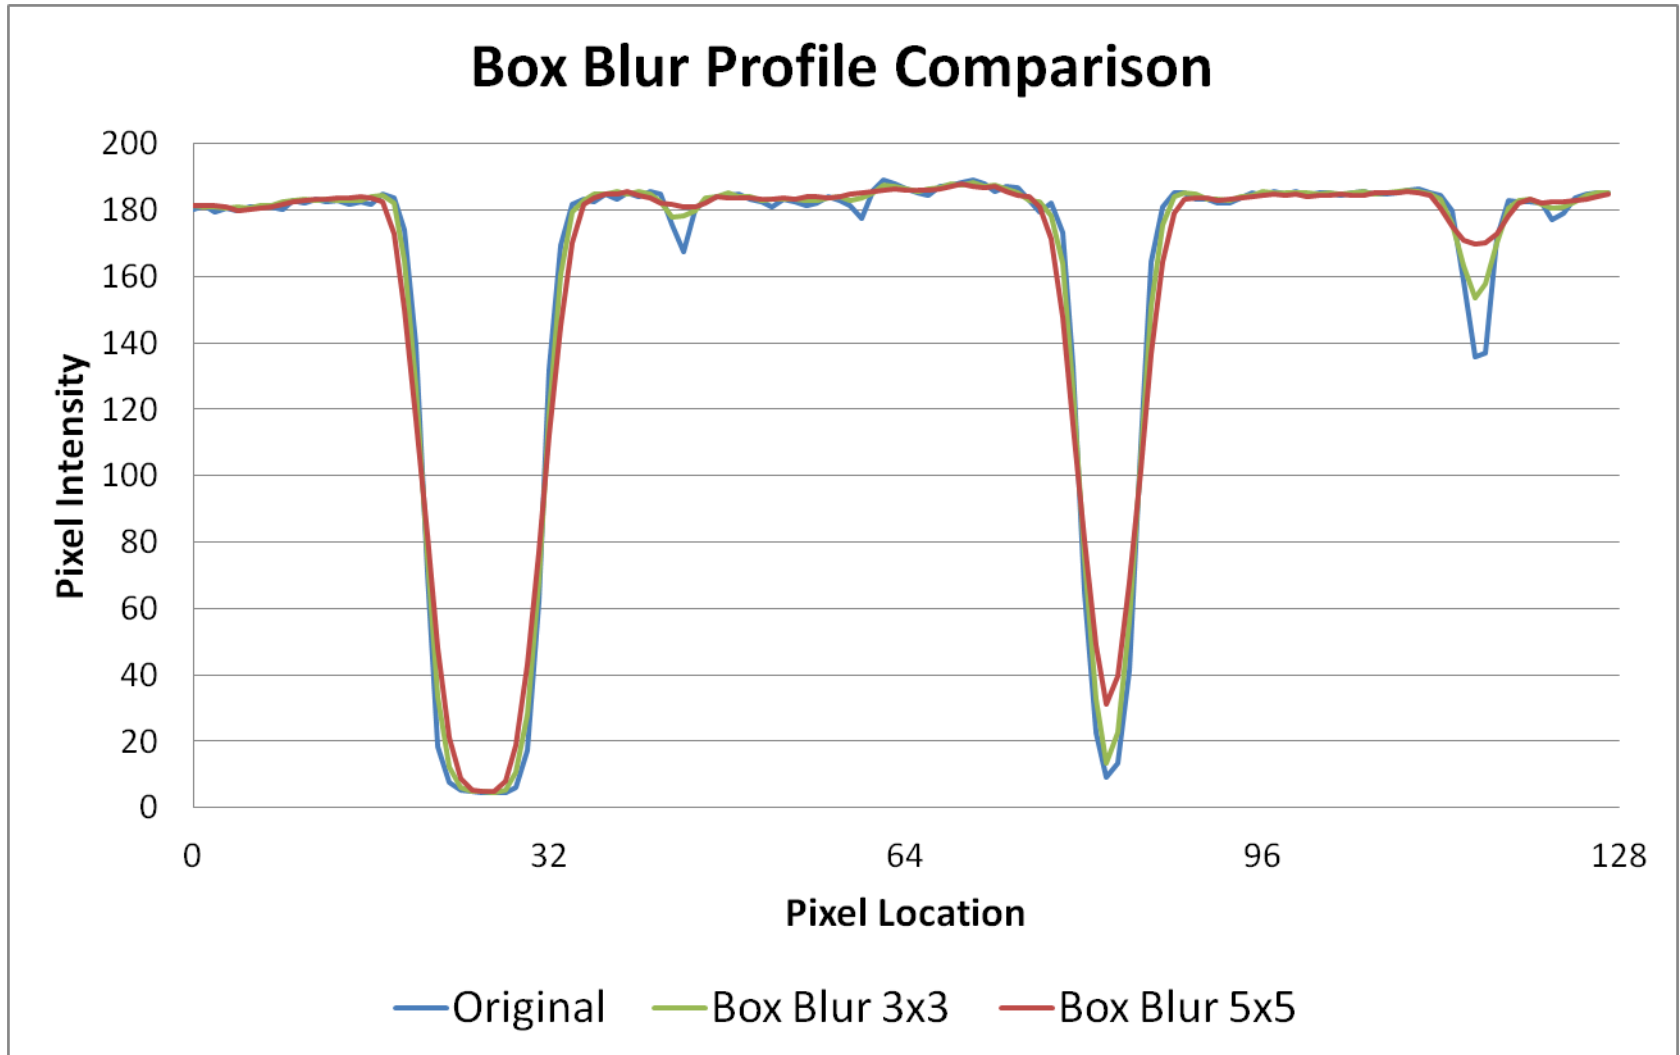


3x3



5x5

# Box Blur Profile Comparison



# Convolution – Sharpen

- Enhances the difference between neighboring pixels
- The greater the difference, the more change in the output pixel

$-1/9$	$-1/9$	$-1/9$
$-1/9$	1	$-1/9$
$-1/9$	$-1/9$	$-1/9$


Sharpen 1

0	$-2/3$	0
$-2/3$	$11/3$	$-2/3$
0	$-2/3$	0


Sharpen 2




# Sharpen Example



Original




Sharpen 1




Sharpen 2

# Color Sharpen Example





Original



Sharpened

# Profile Comparison




# Edge Detection – Sobel Operators

- Two 3×3 kernels are convolved with the original image
- Calculates approximations of the derivatives
  - one for horizontal changes, one for vertical changes
- Final Image produced by taking square root of sum of square of resulting images
- Same method as “Process > Find Edges...”


$-1/2$	$-1$	$-1/2$
$0$	$0$	$0$
$+1/2$	$+1$	$+1/2$

$-1/2$	$0$	$+1/2$
$-1$	$0$	$+1$
$-1/2$	$0$	$+1/2$

# Edge Detection – Examples




# Edge Detection Profile




# Dilation

- Set pixel to the maximum color value within a 3x3 window around the pixel
- Causes objects to grow in size.
- Brightens and fills in small holes

# Dilation Examples – 2x2 Checkerboard




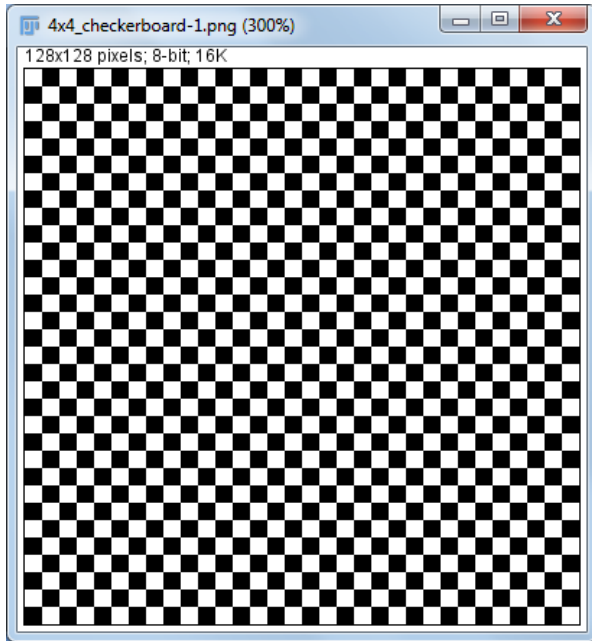
255	255	0
255	255	0
0	0	255



	255	



# Dilation Examples – 4x4 Checkerboard



Only 1 location where 4x4 zeros overlay with 3x3,  
resulting in single pixel of zero


# Erosion

- Set pixel to the minimum color value within a 3x3 window around the pixel
- Causes objects to shrink.
- Darkens and removes small objects

# Erosion Examples – 2x2 Checkerboard

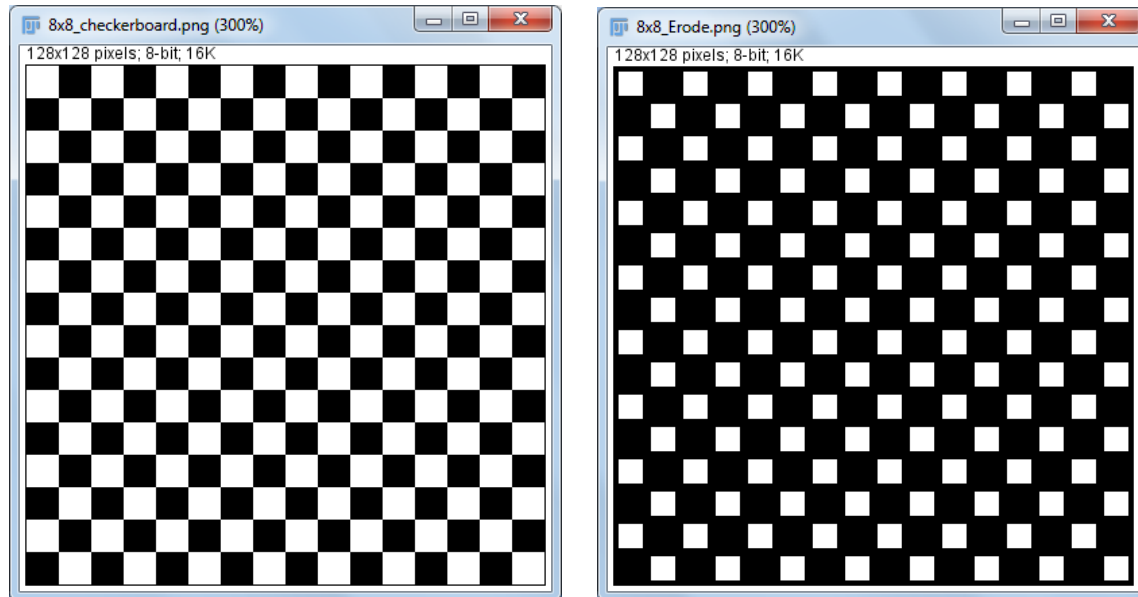


255	255	0
255	255	0
0	0	255



	0	


# Erosion Examples – 8x8 Checkerboard



# Opening (Erosion -> Dilation)



CT Image



Binary Image




Open Result

Note: Using Grays LUT for before and after comparison.


# Close (Dilation -> Erosion)



CT Image



Binary Image




Close Result

Note: Using Grays LUT for before and after comparison.

# Close and Open



CT Image



Binary Image



Close/Open Result


Note: Using Grays LUT for before and after comparison.

# Median Filter

- Replace each pixel in the output image with the median of neighboring pixels in the input image
  - sort 9 input pixel values
  - select middle value (index 4)

[103, 107, 110, 112, 118, 123, 125, 135, 153]



123	107	153
110	125	<u>118</u>
135	112	103




	118	





# Median Filter - Example



# Median Filter – Example 2




# Threshold -> Median Filter



# What do we do with edge pixels?

1. Wrap the image
2. Ignore edge pixels and only compute for those pixels with all neighbors
3. Duplicate edge pixels



# Whole-Image Algorithms


- Histogram Equalization (aka Contrast Stretching)
- Watershed Segmentation
- Rolling Ball Background Subtraction
- ...


# Histogram Equalization

- Increase the global contrast of images
- Intensities are better distributed
- Reveals more detail in images that are over or under exposed



Shift to the right  
implies brighter reds



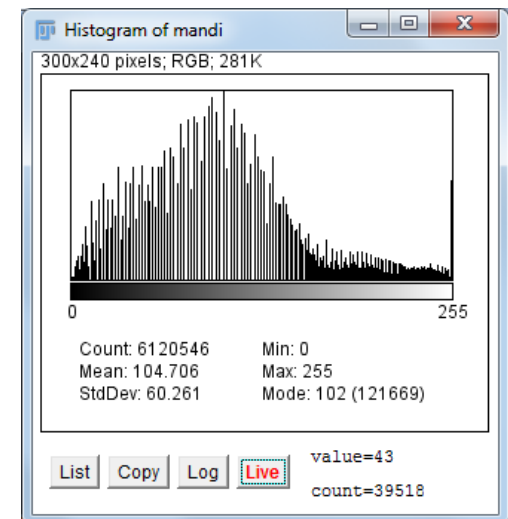
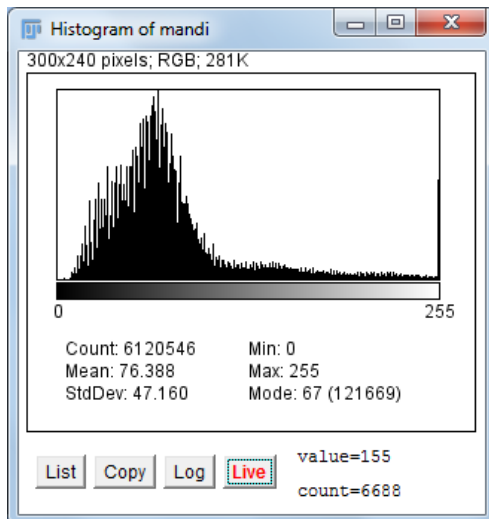
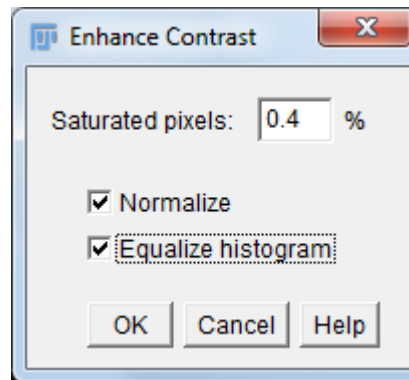





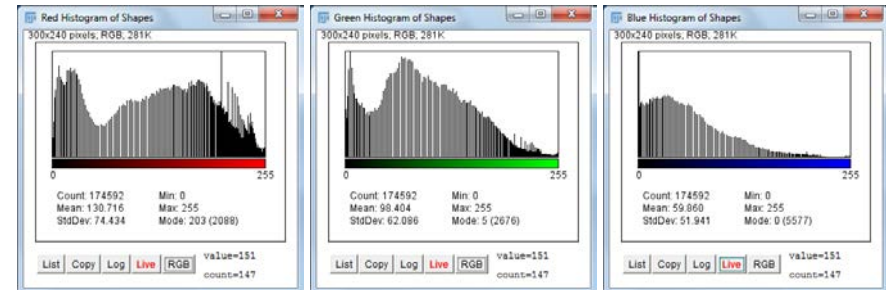
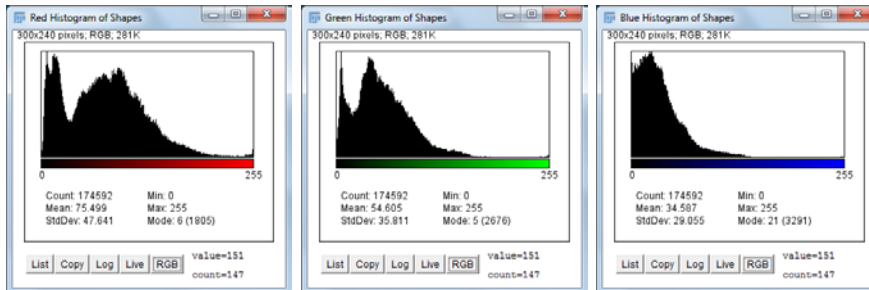
# Histogram Equalization

- Calculate color frequencies - count the number of times each pixel color appear in the image
- Calculate the cumulative distribution function (cdf) for each pixel color – the number of times all smaller color values appear in the image
- Normalize over (0, 255)



# Histogram Equalization Example



# Color Histogram Equalization



# Watershed Segmentation



# Rolling Ball Background Subtraction

- A local background value is determined for every pixel by averaging over a very large ball around the pixel.
- This value is hereafter subtracted from the original image, (hopefully) removing large spatial variations of the background intensities.
- The radius should be set to at least the size of the largest object that is not part of the background.

# Rolling Ball Background Subtraction

